

# Multi-Provider Model Router API Specification

## 1. Overview

The Multi-Provider Model Router is a server application built with FastAPI that routes requests to multiple large language model (LLM) providers based on tenant-specific policies, available models, modalities, and provider health, latency, and cost. It supports multi-tenant authentication, model validation, failover, logging, and metrics monitoring.

## 2. Key Components

- **FastAPI service:** API server exposing endpoints (e.g., `/v1/chat/completions`) for chat completions.
- **Tenant Authentication:** Validates tenants using API keys via `Authorization: Bearer <api_key>` headers.
- **Model Catalog:** YAML-based catalog defining models (e.g., GPT, Claude, Gemini) and their providers, modalities, pricing, and parameters.
- **Routing Policy Engine:** Encapsulates tenant policies that specify primary providers and failover order.
- **Model Router:** Core routing logic that selects the optimal provider considering health, latency, cost, and tenant policy.
- **Provider Adapters:** Abstract connections to different providers, implementing request/response translation.
- **Logging:** Structured logging setup with console, rotating file, and JSON format handlers for production diagnostics.
- **Metrics:** Prometheus integration providing application and routing metrics for usage, latency, failures.
- **Failover:** Automatic retry failover to secondary providers if primary fails.

## 3. API Endpoints

### 3.1 Chat Completions

- **URL:** `/v1/chat/completions`
- **Method:** POST
- **Description:** Accepts chat completion requests, authenticates tenant, validates model, routes to provider, and returns response.
- **Headers:** `Authorization: Bearer <api_key>`
- **Request Body:** JSON containing
  - **model:** String model name (e.g., `gpt-3.5-turbo`)
  - **messages:** List of messages in OpenAI chat format (role + content)
  - **Optional parameters:** `temperature`, `max_tokens`, `top_p`, etc.
- **Response:** JSON response from selected provider routed through the service.

- **Error Codes:**
  - 401 Unauthorized: Missing or invalid API key
  - 403 Forbidden: Tenant not allowed to use requested model
  - 503 Service Unavailable: No primary provider configured
  - 502 Bad Gateway: All providers failed to respond

## 4. Tenant Authentication and Authorization

- Tenants are identified and authenticated via API keys loaded from a YAML config.
- Each tenant's allowed models and providers are defined in tenant configuration.
- Authorization enforced via dependency injection in FastAPI on each request.
- Requests are rejected early if authentication or model authorization fails.

## 5. Models Catalog Specification

- Models and their providers are defined in `models_catalog.yaml`.
- Each model entry includes:
  - `name`: Model identifier (e.g., `gpt-3.5-turbo`)
  - `description`: Human-readable description
  - `categories`: Model classification tags
  - `providers`: List of providers supporting this model
    - Each provider defines:
      - ♣ `name`: Provider identifier (e.g., `openai`)
      - ♣ `base_url`: Provider API base URL
      - ♣ `modalities`: Supported request types (`chat`, `embedding`, `image`, etc.)
      - ♣ `context_length`: Min and max token limits
      - ♣ `prompt_pricing`: Cost per token for different modalities
      - ♣ `supported_parameters`: Request params supported
- Router filters by modality and tenant policies when selecting providers.

## 6. Routing Policy Engine

- Loads tenant-specific routing policies from YAML.
- Policies include:
  - `primary_providers`: Ordered list of preferred providers per tenant.

- failover\_order: List of fallback providers in priority.
- Retrieves tenant policy or defaults to global policy.
- Provides methods to get failover providers excluding the failed one.

## 7. Model Router Logic

- Combines tenant policy, models catalog, and provider health/latency/cost metrics.
- Filters candidate providers by:
  - Tenant primary provider preferences
  - Provider support for requested model and modality
  - Provider health (stub implementation currently always healthy)
- Scores providers by weighted sum of latency and cost.
- Selects the best provider with lowest score.
- Supports failover to secondary providers on request failure.
- Abstracts provider-specific request sending via adapter interface.

## 8. Provider Adapters

- Implement communication logic for each supported provider (OpenAI, Anthropic, Google).
- Translate incoming standardized requests to provider-specific API calls.
- Support sending requests asynchronously and returning parsed responses.

## 9. Logging and Observability

- Structured logging configured with rotating files and JSON formatting.
- Logs include contextual fields (tenant\_id, model\_name, provider, latency, error).
- Middleware logs all incoming requests and outgoing responses with timings and status.
- Provides detailed traceability for routing decisions and failures.

## 10. Metrics and Monitoring

- Prometheus middleware added to FastAPI app to expose /metrics.
- Custom Prometheus metrics include:
  - Request count (model\_router\_requests\_total) by tenant, provider, and model.
  - Request latencies histogram (model\_router\_request\_latency\_seconds).
  - Failure counters (model\_router\_failures\_total).

- Metrics exported for scraping by Prometheus for monitoring and alerting.

## 11. Security Considerations

- API keys securely validated on each request.
- No API secrets exposed externally.
- Sanitize inputs and enforce strict JSON schemas recommended (optional).
- Use HTTPS for all production traffic.
- Optionally implement rate limiting and quotas (quota enforcement currently planned).

## 12. Deployment and Configuration

- Environment-driven configuration supports debug logs, reload, and log level toggles.
- Config files for tenants and models via YAML for easy modifications.
- Providers' API keys and secrets managed securely outside of this catalog.
- Prometheus scraping enabled via `/metrics` endpoint for performance and availability monitoring.

## 13. Future Extensions

- Integrate embedding models and dedicated embedding endpoints.
- Extend adapter plugins for new providers and modalities.
- Enhance health checks with real-time provider status and circuit breakers.
- Add comprehensive unit and integration tests.
- **Implement a database and migrations to maintain the models, tenants without config files to get the more and ease of control.**